

INTRODUCTION TO OPENMP & OPENACC

Kadin Tseng

Boston University

Scientific Computing and Visualization

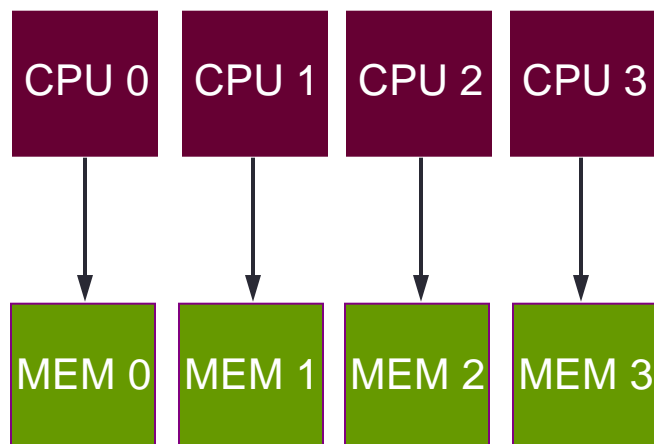
Outline

- Introduction to OpenMP
- Introduction to OpenACC


Introduction to OpenMP (for CPUs)

- Types of parallel machines
 - distributed memory
 - each processor has its own memory address space
 - variable values are independent
 - $x = 2$ on one processor, $x = 3$ on a different processor
 - example: the nodes of a Linux cluster, like SCC's nodes
 - shared memory
 - also called Symmetric Multiprocessing (SMP)
 - single address space for all processors
 - If one processor sets $x = 2$, x will also equal 2 on other processors (unless specified otherwise)
 - example: cores within each SCC node

Shared vs. Distributed Memory

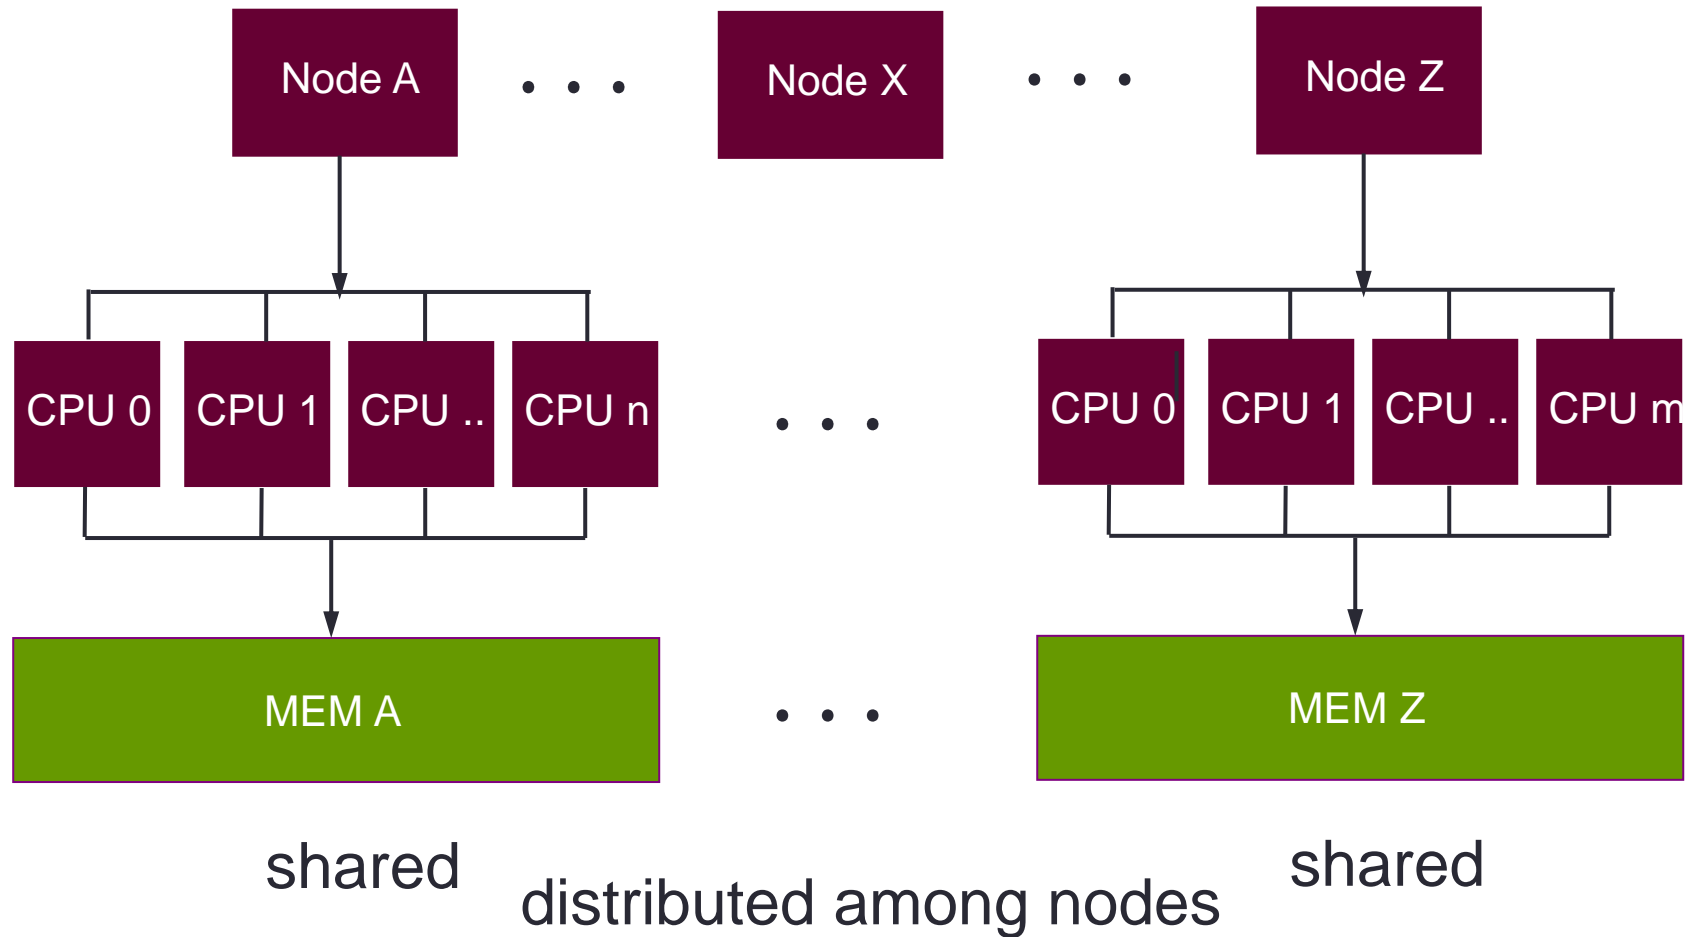


distributed



shared

Shared Computing Cluster (SCC)



The word “shared” in SCC is not for “shared memory.”


What is OpenMP ?

- Application Programming Interface (API) for *multithreaded* parallelization consisting of
 - Source code directives
 - Functions
 - Environment variables
- Advantage
 - Easy to use
 - Incremental parallelization
 - Flexible -- Loop-level or coarse-grain
 - Portable
 - Work on any SMP machine (e.g., each individual SCC node)
- Disadvantage
 - Shared-memory systems only (i.e., not across SCC's nodes)


Basics

- Goal – distribute work among threads
- Two methods to be discussed
 - Loop-level
 - Specified loops are parallelized
 - Used in automatic parallelization tools, like MATLAB PCT
 - Parallel regions
 - Also called “coarse-grained parallelism”
 - Usually used in message-passing (MPI)

Basics (cont'd)



Loop-level



Parallel regions

parallel do & parallel for

- **parallel do** (Fortran) and **parallel for** (C) directives

```
!$omp parallel do
do i = 1, maxi
    c(i) = a(i) + b(i)
enddo
!$omp end parallel do
```

```
#pragma omp parallel for
for(i = 0; i < maxi; i++){
    c[i] = a[i] + b[i];
}
```

↑ Use “c\$” for fixed-format Fortran

- Suppose $\text{maxi} = 1000$ and 4 threads are available
 - Thread 0 gets $i = 1$ to 250
 - Thread 1 gets $i = 251$ to 500
 - Thread 2 gets $i = 501$ to 750
 - Thread 3 gets $i = 751$ to 1000
- Barrier (synchronization) imposed at end of loop

workshare

- For Fortran 90/95 array syntax, the **parallel workshare** directive is analogous to **parallel do**
- Previous example would be:

```
!$omp parallel workshare  
c = a + b  
!$omp end parallel workshare
```

- Also works for **forall** and **where** statements
- No equivalent directive for C/C++

Shared vs. Private

- In parallel region, all variables are *shared* by default
- Loop indices are *private* by default
- What is wrong with the following code segment ?

```
ifirst = 10      ! shared by all threads
!$omp parallel do
do i = 1, imax    ! i is private
    i2 = 2*i      ! i2 is shared
    j(i) = ifirst + i2 ! j also shared
enddo
!$omp end parallel do
```

```
ifirst = 10;      // shared by all threads
#pragma omp parallel for
for(i = 0; i < imax; i++){ // i is private
    i2 = 2*i;          // i2 is shared
    j[i] = ifirst + i2; // j also shared
}
```

Shared vs. Private (cont'd)

Need to declare `i2` with a `private` clause

```
ifirst = 10    !shared by all threads
!$omp parallel do private(i2)
do i = 1, imax ! i is private
    i2 = 2*i    ! i2 different on each thread
    j(i) = ifirst + i2
enddo
!$omp end parallel do
```

```
ifirst = 10;
#pragma omp parallel for private(i2)
for(i = 0; i < imax; i++){ // i is private
    i2 = 2*i; // i2 different on each thread
    j[i] = ifirst + i2;
}
```

Data Dependencies

- Data on one thread can be dependent on data on another thread
- This can result in wrong answers
 - thread 0 may require a variable that is calculated on thread 1
 - answer depends on timing – When thread 0 does the calculation, has thread 1 calculated its value yet?

Data Dependencies (cont'd)

- Example – Fibonacci Sequence 0, 1, 1, 2, 3, 5, 8, 13, ...

```
a(1) = 0
a(2) = 1
do i = 3, 100
    a(i) = a(i-1) + a(i-2)
enddo
```

Lets parallelize on 2 threads.

Thread 0 gets $i = 3$ to 51

Thread 1 gets $i = 52$ to 100

Follow calculation for $i = 52$ on thread 1. What will be values of a at $i - 1$ and $i - 2$?

```
a[1] = 0;
a[2] = 1;
for(i = 3; i <= 100; i++){
    a[i] = a[i-1] + a[i-2];
}
```

More clauses

- Can make **private** the default rather than **shared**
 - Fortran only
 - handy if most of the variables are private
 - can use continuation characters for long lines

```
ifirst = 10
!$omp parallel do      &
!$omp default(private) &
!$omp shared(ifirst,imax,j)
do i = 1, imax
    i2 = 2*i
    j(i) = ifirst + i2
Enddo
!$omp end parallel do
```

More clauses (cont'd)

- Can use **default none**
 - declare all variables (except loop variables) as shared or private
 - If you don't declare any variables, you get a handy list of all variables in loop

More clauses (3)

```
ifirst = 10
!$omp parallel do      &
!$omp default(none)    &
!$omp shared(ifirst,imax,j) private(i2)
do i = 1, imax
    i2 = 2*i
    j(i) = ifirst + i2
enddo
!$omp end parallel do
```

```
ifirst = 10;
#pragma omp parallel for \
                    default(none) \
                    shared(ifirst,imax,j) private(i2)
for(i = 0; i < imax; i++){
    i2 = 2*i;
    j[i] = ifirst + i2;
}
```

Firstprivate

- Suppose we need a running total for each index value on each thread

```
iper = 0
do i = 1, imax
    iper = iper + 1
    j(i) = iper
enddo
```

```
iper = 0;
for(i = 0; i < imax; i++){
    iper = iper + 1;
    j[i] = iper;
}
```

- if **iper** were declared **private**, the initial value would not be carried into the loop

Firstprivate (cont'd)

- Solution – firstprivate clause
- Creates private memory location for each thread
- Copies value from master thread (thread 0) to each memory location

```
iper = 0
!$omp parallel do &
!$omp firstprivate(iper)
do i = 1, imax
    iper = iper + 1
    j(i) = iper
enddo
!$omp end parallel do
```

```
iper = 0;
#pragma omp parallel for \
    firstprivate(iper)
for(i = 0; i < imax; i++){
    iper = iper + 1;
    j[i] = iper;
}
```

Lastprivate

- saves value corresponding to the last loop index
 - "last" in the serial sense

```
!$omp parallel do lastprivate(i)
do i = 1, maxi
  a(i) = b(i)
enddo
a(i) = b(1)
!$omp end parallel do
```

```
#pragma omp parallel for lastprivate(i)
for(i = 0; i < maxi; i++){
  a[i] = b[i];
}
a[i] = b[0];
```

Reduction

- Following example won't parallelize correctly
 - different threads may try to write to **s** simultaneously

```
s = 0.0
!$omp parallel do
do i = 1, maxi
    s = s + a(i)
Enddo
!$omp end parallel do
```

```
s = 0.0;
#pragma omp parallel for
for(i = 0; i < imaxi; i++){
    s = s + a[i];
}
```

Reduction (cont'd)

- Solution is to use the **reduction** clause

```
s = 0.0
!$omp parallel do reduction(+:s)
do i = 1, maxi
    s = s + a(i)
enddo
!$omp end parallel do
```

```
s = 0;
#pragma omp parallel for reduction(+:s)
for(i = 0; i < imaxi; i++){
    s = s + a[i];
}
```

- each thread performs its own reduction (sum, in this case)
- results from all threads are automatically reduced (summed) at the end of the loop

Reduction (3)

- Fortran operators/intrinsics: MAX, MIN, IAND, IOR, IEOR, +, *, -, .AND., .OR., .EQV., .NEQV.
- C operators: +, *, -, /, &, ^, |, &&, ||
- roundoff error may be different than serial case

Conditional Compilation

- For C, C++: conditional compilation performed with `_OPENMP` macro name (defined during compilation with OpenMP turned on*)

```
#ifdef _OPENMP
```

```
... do stuff ...
```

```
#endif
```

- For Fortran: there are two alternatives
 - The above for C works if fortran file named with suffix `.F90` or `.F`
 - Source lines start with `!$` become active with OpenMP turned on*
`!$ print*, 'number of procs =', nprocs`

* How to turn on OpenMP is discussed in [Compile and Run](#) page.

Basic OpenMP Functions

- `omp_get_thread_num()`
 - returns current thread ID; effective inside parallel region
- `omp_set_num_threads(nthreads)`
 - subroutine in Fortran
 - sets number of threads in next parallel region to `nthreads`
 - overrides `OMP_NUM_THREADS` environment variable
 - Effective outside parallel region
- `omp_get_num_threads()`
 - returns number of threads in current parallel region

Some Tips

- OpenMP will do what you tell it to do
 - If you try parallelize a loop with a dependency, it will go ahead and do it!
- Do not parallelize small loops
 - Overhead will be greater than speedup
 - How small is “small”?
 - Answer depends on processor speed and other system-dependent parameters
- Maximize number of operations performed in parallel
 - parallelize outer loops where possible
- For fortran, add “use omp_lib” to include header

Compile and Run on SCC

- Portland Group compilers:
 - Compile with **-mp** flag to turn on OpenMP
 - Depending on the node, can use up to 16 threads
 - `scc1% pgfortran -o myprog myprog.f90 -mp -O3`
 - `scc1% pgcc -o myprog myprog.c -mp -O3`
 - `scc1% setenv OMP_NUM_THREADS 4`
 - `scc1% myprog`
- GNU compilers:
 - Compile with **-fopenmp** flag to turn on OpenMP
 - Depending on the node, can use up to 16 threads
 - `scc1% gfortran -o myprog myprog.f90 -fopenmp -O3`
 - `scc1% gcc -o myprog myprog.c -fopenmp -O3`

Parallel

- **parallel** and **do/for** can be separated into two directives.

```
!$omp parallel do  
do i = 1, maxi  
    a(i) = b(i)  
Enddo  
!$omp end parallel do
```

```
#pragma omp parallel for  
for(i=0; i<maxi; i++){  
    a[i] = b[i];  
}
```

is the same as

```
!$omp parallel  
!$omp do  
do i = 1, maxi  
    a(i) = b(i)  
enddo  
!$omp end parallel
```

```
#pragma omp parallel  
#pragma omp for  
for(i=0; i<maxi; i++){  
    a[i] = b[i];  
}
```

Parallel (cont'd)

- Note that an **end parallel** directive is required.
- Everything within the **parallel** region will run in parallel.
- The **do/for** directive indicates that the loop indices will be distributed among threads rather than duplicating every index on every thread.

Parallel (3)

- Multiple loops in parallel region:

```
!$omp parallel
!$omp do
do i = 1, maxi
    a(i) = b(i)
enddo
!$omp do
do i = 1, maxi
    c(i) = a(2)
enddo
!$omp end parallel
```

```
#pragma omp parallel
#pragma omp for
for(i=0; i<maxi; i++){
    a[i] = b[i];
}
#pragma omp for
for(i=0; i<maxi; i++){
    c[i] = a[2];
}
#pragma omp end parallel
```

- parallel** directive has a significant overhead associated with it.
- The above example has the potential to be faster than using two **parallel do/parallel for** directives.

Coarse-Grained

- OpenMP is not restricted to loop-level, or fine-grained, parallelism.
- The **!\$omp parallel** or **#pragma omp parallel** directive duplicates subsequent code on all threads until a **!\$omp end parallel** or **#pragma omp end parallel** directive is encountered.
- Allows parallelization similar to “MPI paradigm.”

Coarse-Grained (cont'd)

```
!$omp parallel &  
!$omp private(myid,istart,iend,nthreads,nper)  
nthreads = omp_get_num_threads()  
nper = imax/nthreads  
myid = omp_get_thread_num()  
istart = myid*nper + 1  
iend = istart + nper - 1  
call do_work(istart,iend)  
do i = istart, iend  
    c(i) = a(i)*b(i) + ...  
enddo  
!$omp end parallel
```

```
#pragma omp parallel \  
#pragma omp private(myid,istart,iend,nthreads,nper)  
nthreads = OMP_GET_NUM_THREADS();  
nper = imax/nthreads;  
myid = OMP_GET_THREAD_NUM();  
istart = myid*nper;  
iend = istart + nper - 1;  
do_work(istart,iend);  
for(i=istart; i<=iend; i++){  
    c[i] = a[i]*b[i] + ...  
}
```


Thread Control Directives

Barrier

- **barrier** synchronizes threads

```
!$omp parallel private(myid,istart,iend)
call myrange(myid,istart,iend)
do i = istart, iend
    a(i) = a(i) - b(i)
enddo
!$omp barrier
myval(myid+1) = a(istart) + a(1)
!$omp end parallel
```

```
#pragma omp parallel private(myid,istart,iend)
myrange(myid,&istart,&iend);
for(i=istart; i<=iend; i++){
    a[i] = a[i] - b[i];
}
#pragma omp barrier
myval[myid] = a[istart] + a[0]
```

- Here **barrier** assures that a(1) or a[0] is available before computing **myval**

Master

- if you want part of code to be executed only on master thread, use **master** directive
- “non-master” threads will skip over **master** region and continue

Master Example - Fortran

```
!$OMP PARALLEL PRIVATE(myid,istart,iend)
  call myrange(myid,istart,iend)
  do i = istart, iend
    a(i) = a(i) - b(i)
  enddo
!$OMP BARRIER
!$OMP MASTER
write(21) a
!$OMP END MASTER
call do_work(istart,iend)
!$OMP END PARALLEL
```

Master Example - C

```
#pragma omp parallel private(myid,istart,iend)
  myrange(myid,&istart,&iend);
  for(i=istart; i<=iend; i++){
    a[i] = a[i] - b[i];
  }
#pragma omp barrier
#pragma omp master
  fwrite(fid,sizeof(float),iend-istart+1,a);
#pragma omp end master
do_work(istart,iend);
#pragma omp end parallel
```

Single

If you :

- want part of code to be executed only by a single thread
- don't care whether or not it's the master thread

The use **single** directive

- Unlike the **end master** directive, **end single** has barrier

Single Example - Fortran

```
!$OMP PARALLEL PRIVATE(myid,istart,iend)
  call myrange(myid,istart,iend)
  do i = istart, iend
    a(i) = a(i) - b(i)
  enddo
!$OMP BARRIER
!$OMP SINGLE
write(21) a
!$OMP END SINGLE
call do_work(istart,iend)
!$OMP END PARALLEL
```

Single Example - C

```
#pragma omp parallel private(myid,istart,iend)
myrange(myid,istart,iend);
for(i=istart; i<=iend; i++){
    a[i] = a[i] - b[i];
}
#pragma omp barrier
#pragma omp single
fwrite(fid,sizeof(float),nvals,a);
#pragma omp end single
do_work(istart,iend);
```


Critical

If you have code section that:

1. must be executed by every thread
2. threads may execute in any order
3. threads must not execute simultaneously

This does not have a barrier.

Critical Example - Fortran

```
!$OMP PARALLEL PRIVATE(myid,istart,iend)
  call myrange(myid,istart,iend)
  do i = istart, iend
    a(i) = a(i) - b(i)
  enddo
  !$OMP CRITICAL
  call mycrit(myid,a)
  !$OMP END CRITICAL
  call do_work(istart,iend)
!$OMP END PARALLEL
```

Critical Example - C

```
#pragma omp parallel private(myid,istart,iend)
myrange(myid,istart,iend);
for(i=istart; i<=iend; i++){
    a[i] = a[i] - b[i];
}
#pragma omp critical
mycrit(myid,a);
#pragma omp end critical
do_work(istart,iend);
#pragma omp end parallel
```

Ordered

- Suppose you want to write values in a loop:

```
do i = 1, nproc  
    call do_lots_of_work(result(i))  
    write(21,101) i, result(i)  
enddo
```

```
for(i = 0; i < nproc; i++){  
    do_lots_of_work(result[i]);  
    fprintf(fid,"%d %f\n",i,result[i]);  
}
```

- If loop were parallelized, could write out of order
- **ordered** directive forces serial order

Ordered (cont'd)

```
!$omp parallel do
do i = 1, nproc
  call do_lots_of_work(result(i))
  !$omp ordered
  write(21,101) i, result(i)
  !$omp end ordered
enddo
```

```
#pragma omp parallel for
for(i = 0; i < nproc; i++){
  do_lots_of_work(result[i]);
  #pragma omp ordered
  fprintf(fid,"%d %f\n",i,result[i]);
  #pragma omp end ordered
}
```

- Since `do_lots_of_work` takes a lot of time, most parallel benefit will be realized

Schedule

- **schedule** refers to the way in which loop indices are distributed among threads
- **([static[, chunk]])**
 - static is the default
 - each thread is assigned a contiguous chunk of indices in thread number order
 - number of indices assigned to each thread is as equal as possible
 - Chunk size may be specified
- **(dynamic[, chunk])**
 - Good way for varying work load among loop iterations

Hands On Exercise

- Parallelize a serial C or Fortran code with OpenMP
- The code invokes a function multiple times via a for/do loop.
Parallelize the loop with an OpenMP directive
- The function, *mywork*, also has a loop (no need to parallelize) whose iteration count is randomly chosen each time it is invoked. Hence, *mywork*'s workload varies with each iteration. This causes load imbalance. Use an appropriate OpenMP clause to address this problem
- A serial *myprog.c* and *myprog.f90* are available for you to start with. Look in /scratch/kadin

Introduction to OpenACC

- OpenMP is for CPUs, OpenACC is for GPUs
- Has runtime library like OpenMP
- Can mix OpenMP with OpenACC

Laplace Equation

$$\frac{\partial^2 u}{\partial x^2} + \frac{\partial^2 u}{\partial y^2} = 0$$

Boundary Conditions:

$$u(x, 0) = 0 \quad 0 \leq x \leq 1$$

$$u(x, 1) = 0 \quad 0 \leq x \leq 1$$

$$u(0, y) = u(1, y) = 0 \quad 0 \leq y \leq 1$$

Finite Difference Numerical Discretization

Discretize equation by centered-difference yields:

$$u_{i,j}^{n+1} \cong \frac{u_{i+1,j}^n + u_{i-1,j}^n + u_{i,j+1}^n + u_{i,j-1}^n}{4} \quad i = 1, 2, \dots, m; \quad j = 1, 2, \dots, m$$


where n and $n+1$ denote the current and the next time step, respectively, while


$$\begin{aligned} u_{i,j}^n &= u^n(x_i, y_j) \quad i = 0, 1, 2, \dots, m+1; \quad j = 0, 1, 2, \dots, m+1 \\ &= u^n(i\Delta x, j\Delta y) \end{aligned}$$

For simplicity, we take

$$\Delta x = \Delta y = \frac{1}{m+1}$$

Computational Domain





Laplace Solver with OpenMP

```
!$omp parallel do shared(m, n, up, u) reduction( max:error )  
  do j=1,m  
    do i=1,n  
       $up(i,j) = ( u(i+1, j) + u(i-1, j) + u(i, j-1) + u(i, j+1) ) * 0.25$   
       $error = \max( error, \text{abs}(up(i, j) - u(i, j)) )$   
    end do  
  end do  
!$omp end parallel do
```

Corresponding C parallel directive is:

```
#pragma parallel for shared(m,n,up,u) reduction( max:error )
```

Laplace Solver with OpenACC

!\$acc kernels

```
do j=1,m
  do i=1,n
    up(i, j) = ( u(i+1, j) + u(i-1, j) + u(i, j-1) + u(i, j+1) ) * 0.25
    error = max( error, abs(up(i, j) - u(i, j)) )
  end do
end do
```

!\$acc end kernels

- **#pragma acc kernels** for C
- Alternatively, **!\$acc parallel loop**, **!\$acc parallel** and **!\$acc loop** are available. Good to start with **kernels** . . .

OpenACC data clause

```
#pragma acc data copy(u), create(up)
```

```
while ( error > tol && iter < iter_max ) { error = 0.0;
```

```
#pragma acc kernels
```

```
    for (int i = 1; i <= n; i++) {
```

```
        for (int j = 1; j <= m; j++ ) {
```

```
            up[i][j] = ( u[i][j+1] + u[i][j-1] + u[i-1][j] + u[i+1][j]) * 0.25;
```

```
            error = fmax( error, fabs(up[i][j] - u[i][j])); }
```

```
        }
```

```
#pragma acc kernels
```

```
    for (int i = 1; i <= n; i++) {
```

```
        for (int j = 1; j <= m; j++ ) {
```

```
            u[i][j] = up[i][j]; }
```

```
        }
```

```
    iter++;
```

```
}
```

- *copy into and out of region*
- *copyin only on in*
- *copyout only on out*
- *create within region*
- *Default is copy without data*

OpenACC on SCC

- Hardware (GPU)
 - Each node has 3 Nvidia Tesla M2050 GPUs – Nehalem class **buy-in** 12-core nodes
 - 3 GB memory/gpu, 448 cores/gpu
 - Each node has 8 Nvidia Tesla M2070 GPUs – Nehalem class **public** 12-core nodes
 - 6 GB memory/gpu, 448 cores/gpu
- Compiler
 - On the SCC, only Portland Group compilers support OpenACC
 - Current (default) version is 13.5
- How to compile codes with OpenACC directives
 - `scc1% pgfortran -o prog prog.f90 -tp=nehalem -acc -ta=nvidia,time -Minfo=accel`
 - `scc1% pgcc -o myprog myprog.c -tp=nehalem -acc -ta=nvidia,time -Minfo=accel`
 - `-tp=nehalem` below creates executable for Intel Nehalem class
 - `-acc` engages the OpenACC API
 - `-ta=nvidia,time` links with Nvidia library for timing data in accelerator region
 - `-Minfo=accel` instructs compiler to display warning and error messages
- Tips from PGI
 - http://www.pgroup.com/resources/openacc_tips_fortran.htm

OpenACC on SCC (cont'd)

- How to run jobs

Login nodes have no GPUs. Must run via batch scheduler

- **Interactive batch** -- for program development and debugging

Example: 1 gpu, 1 cpu, 4 hours of estimated runtime

- `scc1% qsh -l gpus=1 -l h_rt=04:00:00`
- `-l gpus=G/C`; G = number of GPUs, C = number of CPU cores

- **Background Batch** -- for production runs

Example: 8 GPUs, 12 CPUs, 4 hours of runtime

- `scc1% qsub -l gpus=0.667 -pe omp 12 -l h_rt=04:00:00`
- `-l gpus = G/C = 8/12 = 0.667`
- `scc1% qsub mymscript` (*mymscript* includes above parameters)